

Lecture 2: Parallel computing

Informatik elective: GPU Computing

Pratik Nayak

Licensed under 

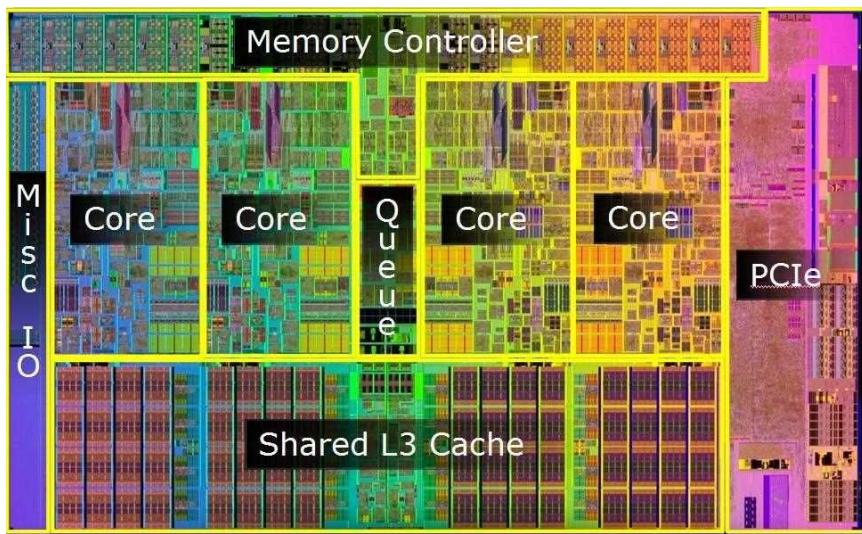


In this session

- Execution strategies
- Task graphs: critical path, concurrency, task decompositions
- Estimating work, memory and cache complexities
- NUMA and cache effects
- Estimating performance
- Roofline model, arithmetic intensity, machine balance
- Scaling: Strong and weak.

Recall terminology: Thread and Core

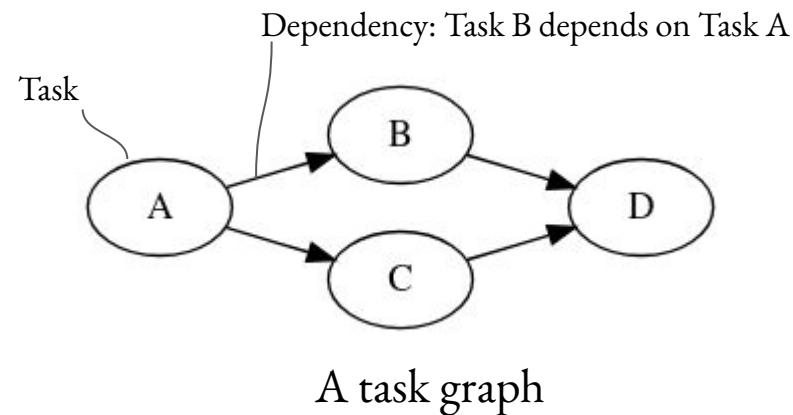
- Thread (software-level): “thread of execution”: an ordered sequence of instructions (software)
- Core (hardware-level): One processor within a CPU die (hardware).
- Multi-core (hardware-level): Multiple processors capable of independent execution within one CPU die.



[Intel multi-core CPU architecture, Intel Corp]

Basic definitions

- Task: Unit of execution or a unit of work
- Dependency: Requirement that enforces a partial ordering of tasks
- Computing Resource: A processing unit on which a task can be executed.



Execution strategies

Two tasks, A and B need to be completed.

Questions:

1. What are the dependencies ?
2. What execution resources are available ?

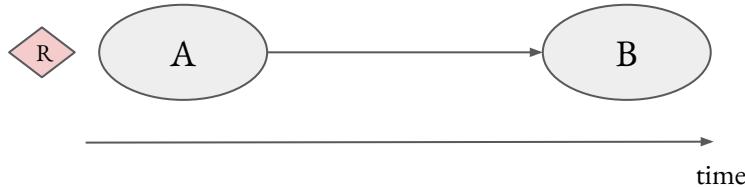
Execution strategies

Two tasks, A and B need to be completed.

Questions:

1. What are the dependencies ?
2. What execution resources are available ?

- Serial execution: Execute task A and then execute task B



- If B depends on A
- OR
- If only one compute resource is available

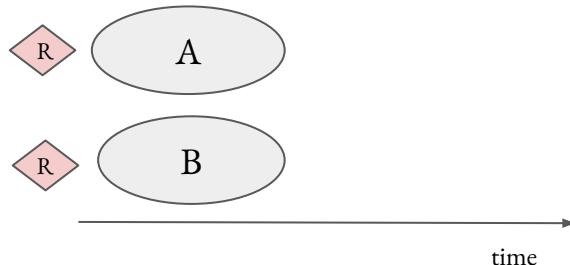
Execution strategies

Two tasks, A and B need to be completed.

Questions:

1. What are the dependencies ?
2. What execution resources are available ?

- Parallel execution: Execute A and B at the same time on separate resources



- A and B are independent AND
- Multiple resources are available

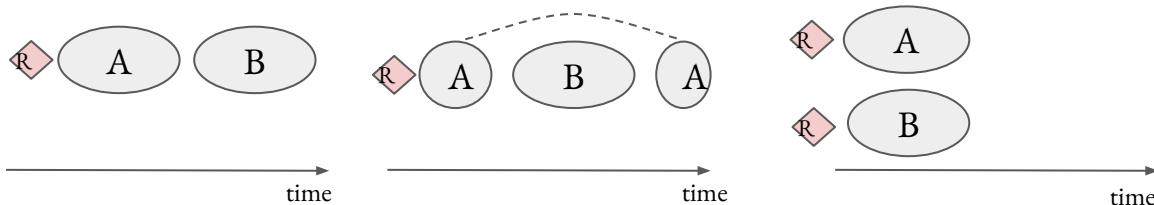
Execution strategies

Two tasks, A and B need to be completed.

Questions:

1. What are the dependencies ?
2. What execution resources are available ?

- Concurrent execution: Execute A and execute B at the same time (no specification on resources)



→ A and B are independent

Execution strategies

Two tasks, A and B need to be completed.

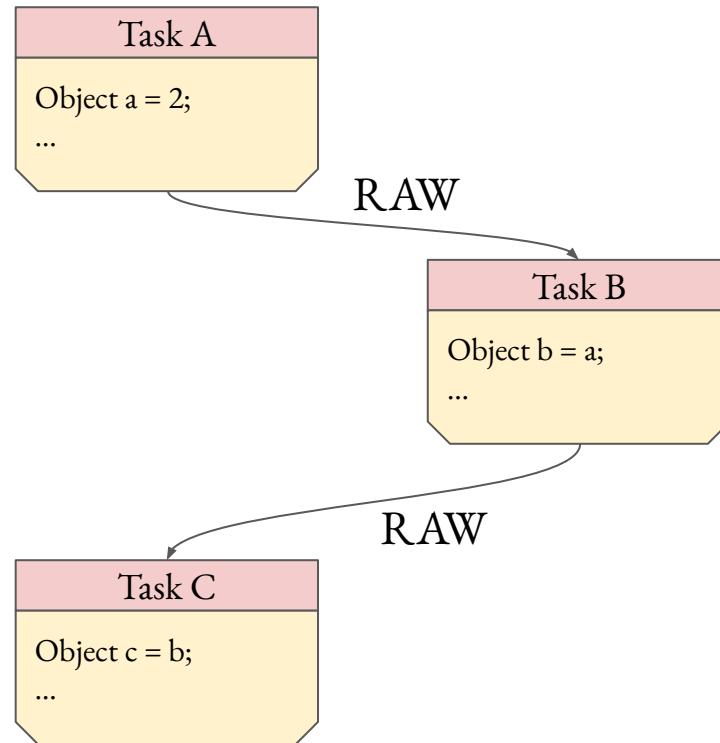
Questions:

1. What are the dependencies ?
2. What execution resources are available ?

- *Serial execution:* Execute task A and then execute task B
- *Concurrent execution:* Execute A and execute B at the same time (no restriction on resources)
- *Parallel execution:* Execute A and B at the same time on separate resources

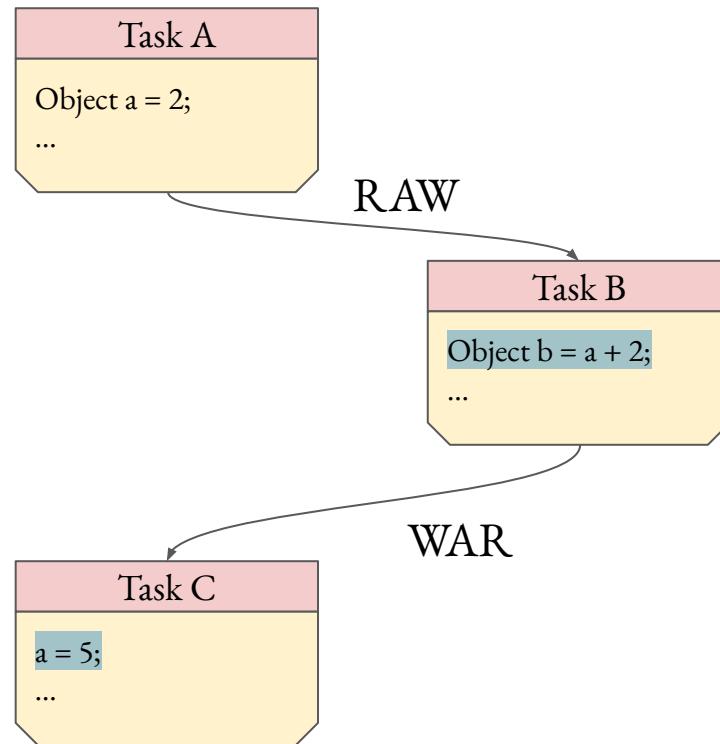
Data dependencies

- (*RAW*) Read-after-write dependency: True dependency



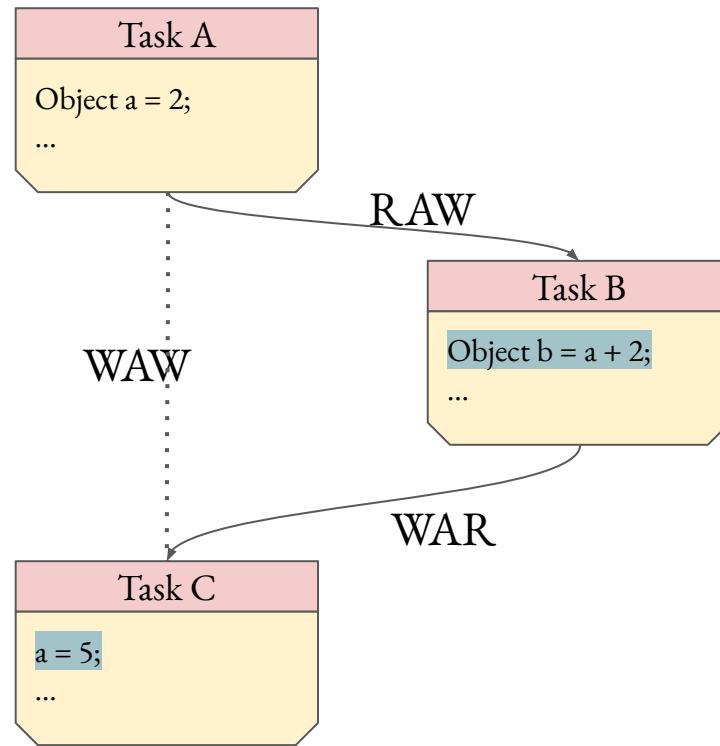
Data dependencies

- (*RAW*) Read-after-write
dependency: True dependency
- (*WAR*) Write-after-read
dependency: Anti-dependency



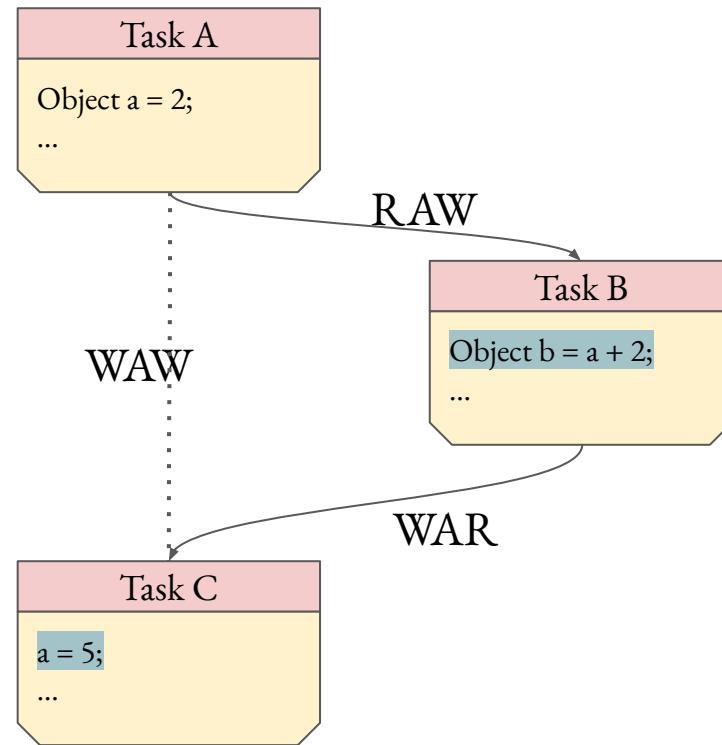
Data dependencies

- (*RAW*) Read-after-write
dependency: True dependency
- (*WAR*) Write-after-read
dependency: Anti-dependency
- (*WAW*) Write-after-write
dependency: Output dependency



Data dependencies

- (*RAW*) Read-after-write
dependency: True dependency
- (*WAR*) Write-after-read
dependency: Anti-dependency
- (*WAW*) Write-after-write
dependency: Output dependency

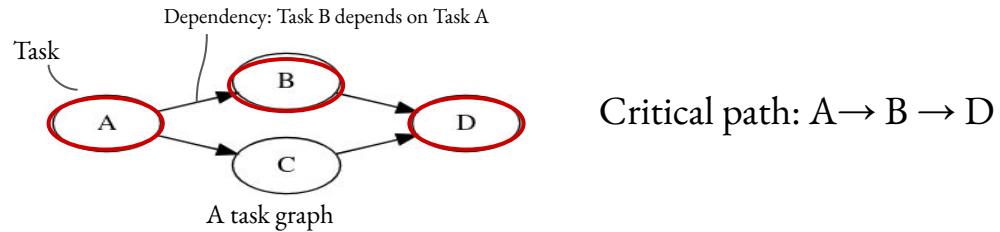


Race conditions and deadlocks

- Race condition: Undefined behaviour based on different thread execution order
 - Data races:
 - Two threads writing to same location (write-write conflict)
 - One thread reads, while other writes to same location (read-write conflict)
 - Incorrect interleaving:
 - Intermediate state exposed to other threads
- Deadlocks:
 - A dependency cycle between concurrent threads:
 - Thread t_1 depends on t_2 and t_2 depends on t_1 . No forward progress.

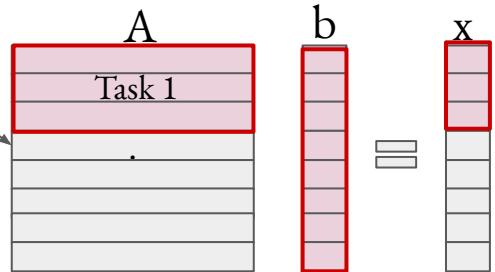
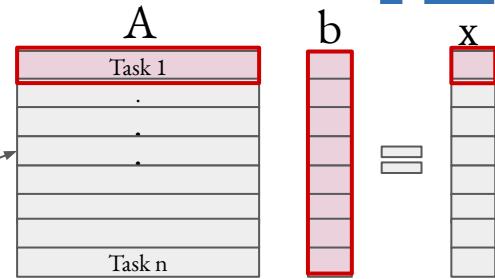
Decomposition and tasking

- Given some computation, to enable parallel processing, we need to:
 - Decompose the computations into smaller tasks
 - Assign the tasks to available resources.
- Decomposition into tasks is generally represented in terms of a task-dependency graph
 - A directed graph with nodes representing the tasks, and edges representing the dependencies.
 - Directed graph → Represents a sequence of tasks
 - Critical path: Longest directed path between a start node and a finish node of a task dependency graph



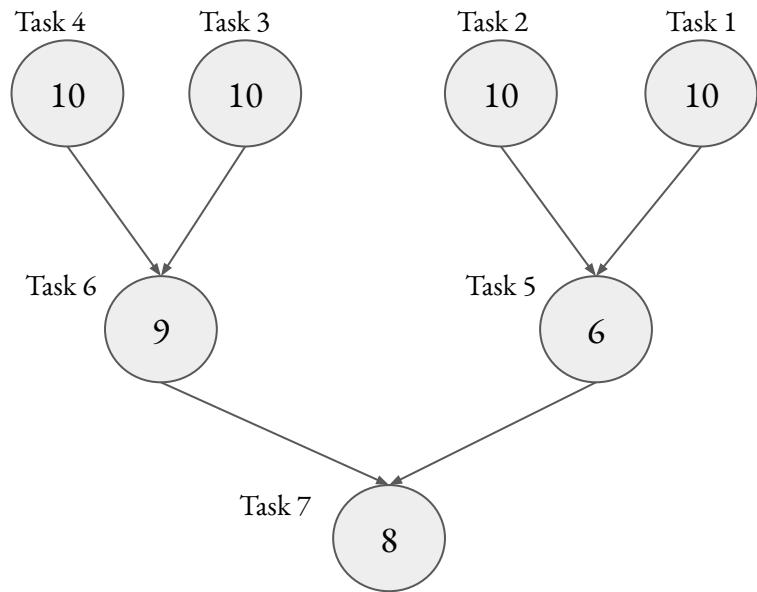
Decomposition and tasking

- Choice of decompositions → Task granularity (Gr):
 - Large number of small tasks (fine-grained decomposition)
 - Small number of large tasks (coarse-grained decomposition)
- Degree of concurrency (#C) → Number of tasks that can execute in parallel
 - Maximum #C → Largest number of tasks at any point of execution
 - Average #C → Average number of tasks that can be executed concurrently → total work/critical path length



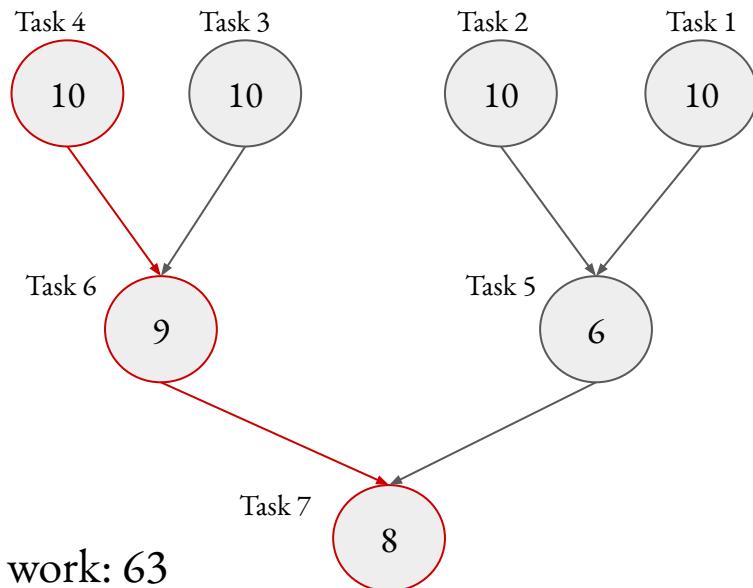
$$\#C \propto \frac{1}{Gr}$$

Critical path: Examples



[Introduction to Parallel Computing, Grama, Gupta, Karypis]

Critical path: Examples



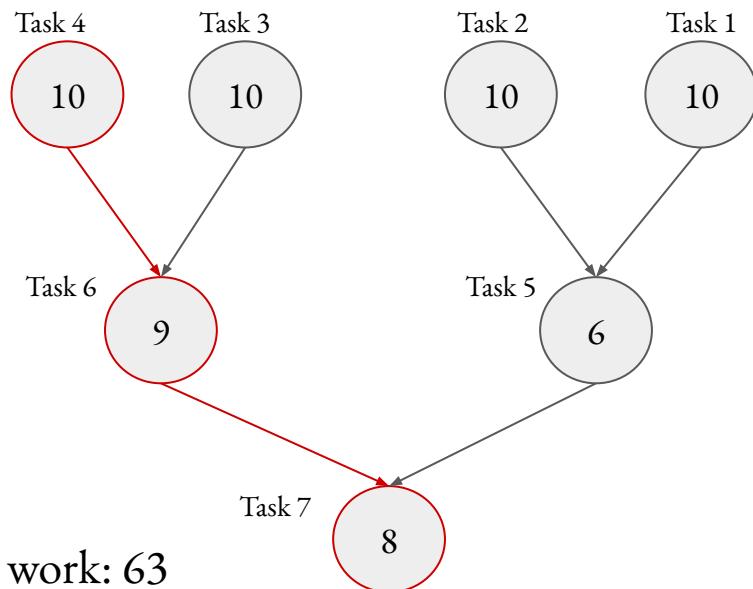
Total work: 63

Critical path length: 27

Average concurrency ~ 2.3

[Introduction to Parallel Computing, Grama, Gupta, Karypis]

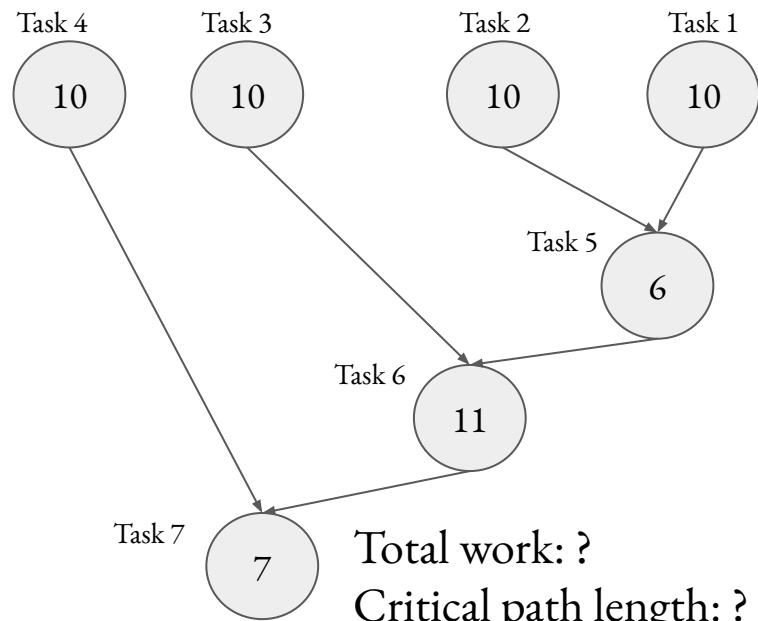
Critical path: Examples



Total work: 63

Critical path length: 27

Average concurrency ~ 2.3



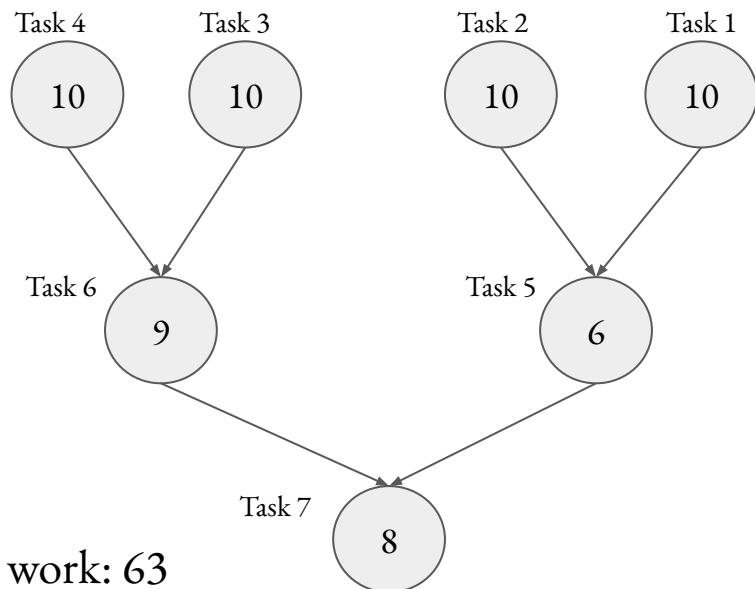
Total work: ?

Critical path length: ?

Average concurrency: ?

[Introduction to Parallel Computing, Grama, Gupta, Karypis]

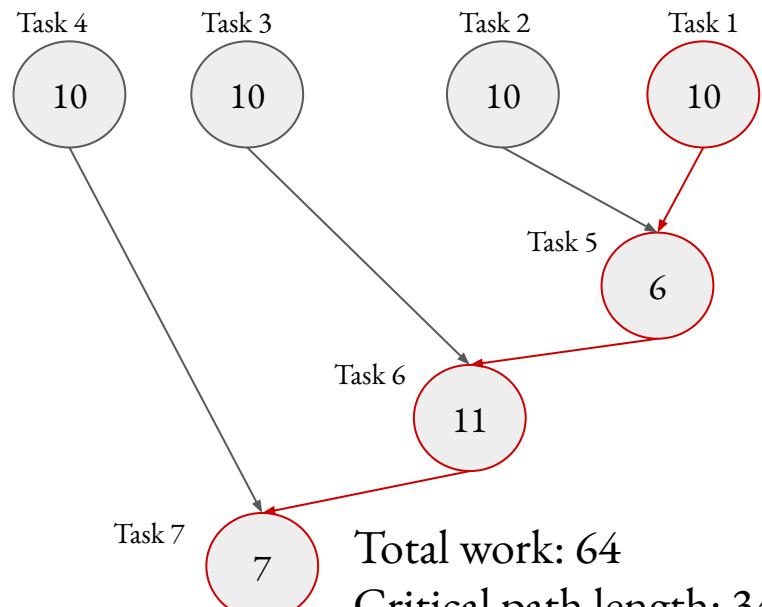
Critical path: Examples



Total work: 63

Critical path length: 27

Average concurrency ~ 2.3



Total work: 64

Critical path length: 34

Average concurrency ~ 1.9

[Introduction to Parallel Computing, Grama, Gupta, Karypis]

Algorithmic complexity

- Assume T_1 to be the amount of work and p be the number of available parallel resources
- Given a critical path length (span) of T_∞ , the lower bound for algorithmic complexity is given by

$$T_p \geq \max\left(\frac{T_1}{p}, T_\infty\right)$$

- A greedy scheduler (no unenforced idleness), for example achieves (upper-bound)

$$T_{p,greedy} \leq \frac{T_1}{p} + T_\infty$$

[Brent 1975 and Introduction to Parallel Computing, Grama, Gupta, Karypis]

Work and memory complexities

- (Work complexity): Given some computation, parameterized by an input size n , the work complexity is defined by $W(n)$: The total number of basic operations (add, mul, divide etc) required.
 - Example (vector addition): Addition of two vectors: u and v , each of size n has a work complexity of $W(n) = n$
 - Example (dot product): Dot product of two vectors: u and v , each of size n has a work complexity of $W(n) = ?$

Work and memory complexities

- (Work complexity): Given some computation, parameterized by an input size n , the work complexity is defined by $W(n)$: The total number of basic operations (add, mul, divide etc) required.
 - Example (vector addition): Addition of two vectors: u and v , each of size n has a work complexity of $W(n) = n$
 - Example (dot product): Dot product of two vectors: u and v , each of size n has a work complexity of $W(n) = 2n$

Work and memory complexities

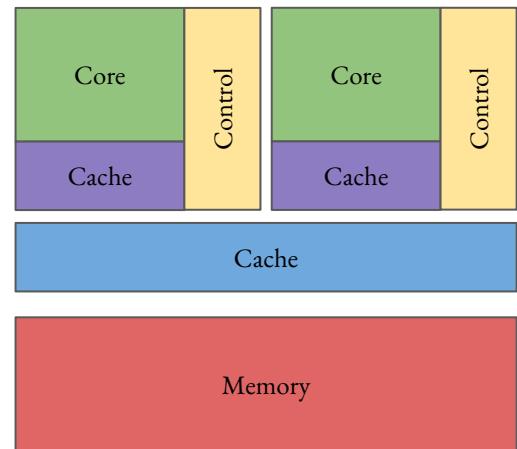
- (Work complexity): Given some computation, parameterized by an input size n , the work complexity is defined by $W(n)$: The total number of basic operations (add, mul, divide etc) required.
- (Memory complexity): Given some computation, parameterized by an input size n , the memory complexity is defined by $M(n)$: The total number of memory operations (reads and writes) required.
 - Ex (vector addition): Addition of two vectors: u and v , each of size n has a memory complexity of $M(n) = 3n$
 - Ex (dot product): Dot product of two vectors: u and v , each of size n has a memory complexity of $M(n) = ?$

Work and memory complexities

- (Work complexity): Given some computation, parameterized by an input size n , the work complexity is defined by $W(n)$: The total number of basic operations (add, mul, divide etc) required.
- (Memory complexity): Given some computation, parameterized by an input size n , the memory complexity is defined by $M(n)$: The total number of memory operations (reads and writes) required.
 - Ex (vector addition): Addition of two vectors: u and v , each of size n has a memory complexity of $M(n) = 3n$
 - Ex (dot product): Dot product of two vectors: u and v , each of size n has a memory complexity of $M(n) = 2n + 1$

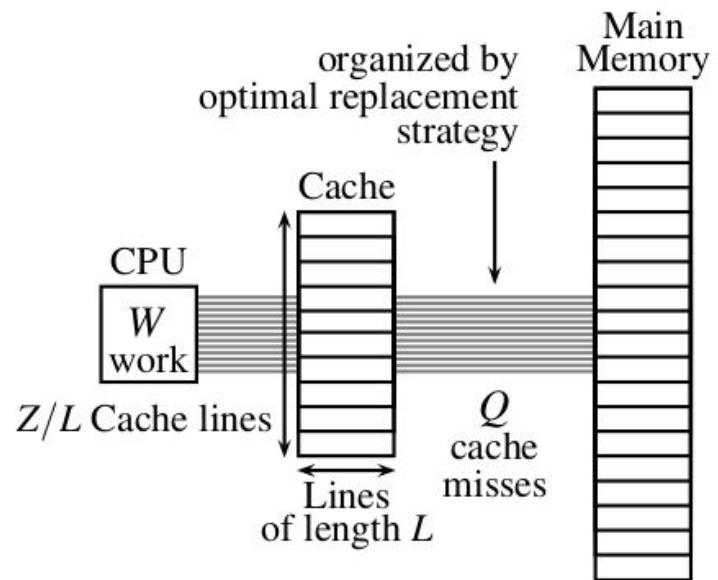
Caching

- Intermediate memory storage between main memory and compute
- Compute unit looks for data in the closest cache:
 - If data is found → cache hit. load data from cache
 - Inability to find data → cache miss, (#cache misses → Q)
 - Recursively look for data in higher-level caches
- Fetches from cache always in blocks called cache lines (L)
- Fetches are hierarchical:
 - Data must be hierarchically fetched into each level of cache, if it needs to be fetched from main memory.



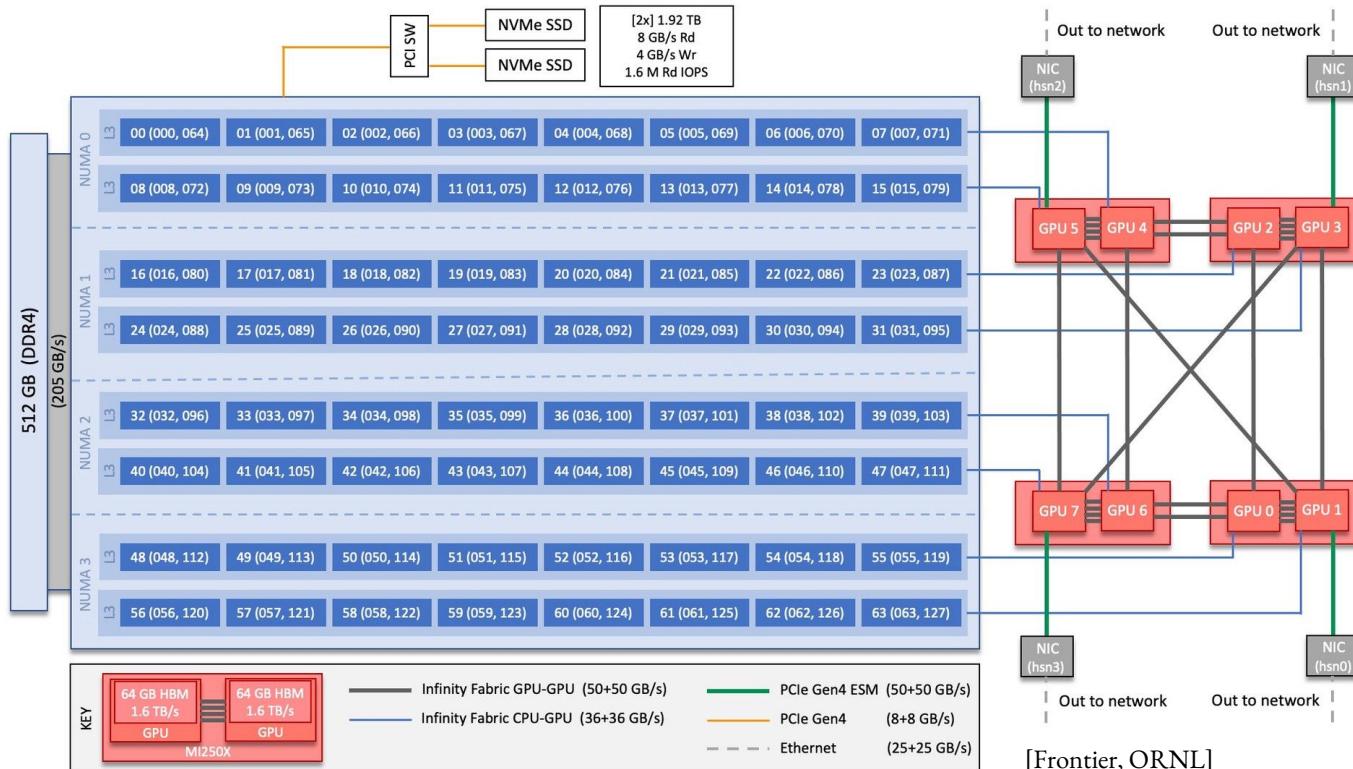
Cache complexity

- Ideal cache model:
 - L : Length of Cache line
 - $Z = \Omega(L^2)$: Cache size (tall cache)
- Cache complexity of an algorithm:
 - Input size n
 - Number of cache misses: $Q(n; Z, L)$
- Cache aware v/s cache oblivious algorithms
 - Cache aware: Q parameterized on compile-time/runtime parameters
 - Cache-oblivious: Q independent of compile-time/runtime parameters



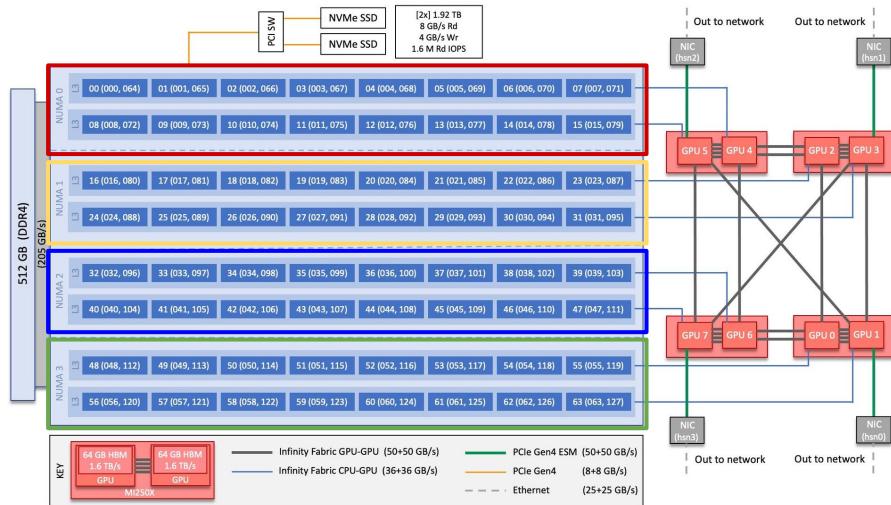
[Frigo et.al , Cache-Oblivious Algorithms, 2012]

Multi-core systems and NUMA



Multi-core systems and NUMA

- 4 NUMA (Non-Uniform Memory Access) domains per node
 - 2 L3 caches per NUMA domain
- Mapping compute region to appropriate memory region crucial.
- GPUs only connected to some NUMA domains



[Frontier, ORNL]

Measuring compute performance: FLOP/s

$$FLOP/s = \text{racks} \times \frac{\text{nodes}}{\text{rack}} \times \frac{\text{sockets}}{\text{node}} \times \frac{\text{cores}}{\text{socket}} \times \frac{\text{cycles}}{\text{second}} \times \frac{FLOP}{\text{cycle}}$$

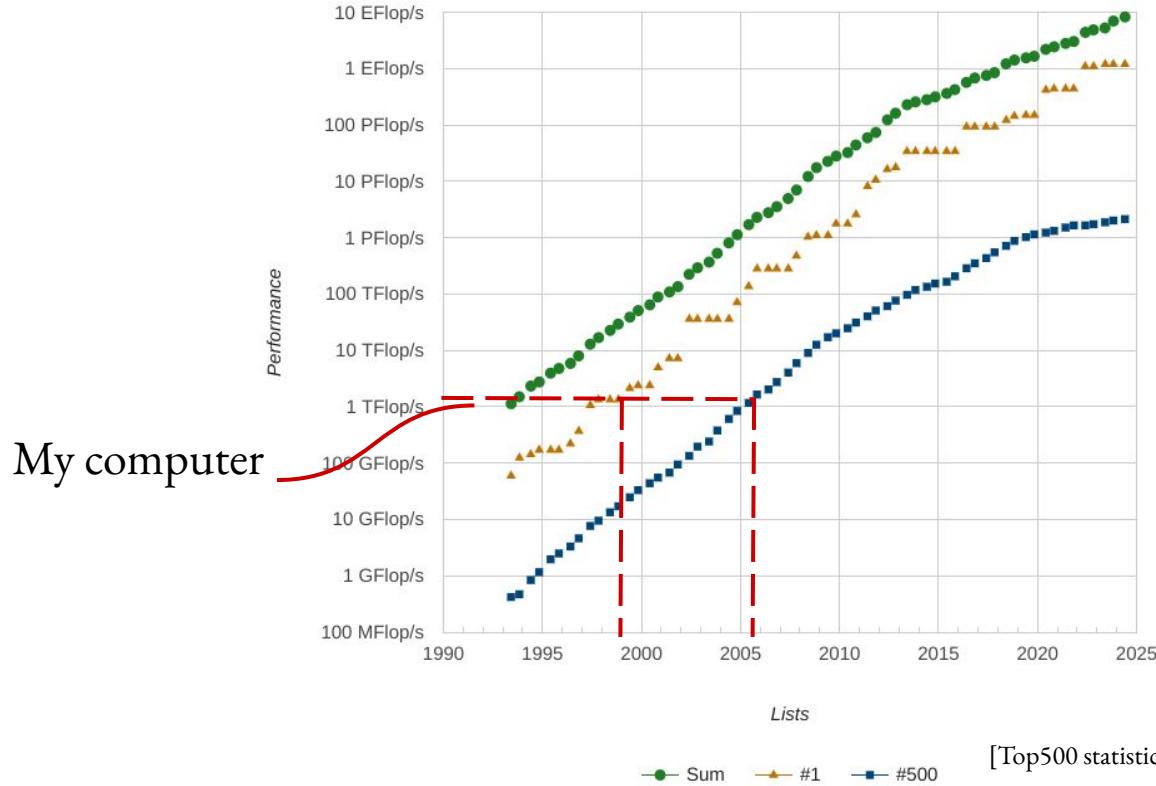
- My computer: 32 FLOP/cycle, 6 GHz, 8 cores/socket, 1 socket: $1536 \text{ GFLOP/s} = 1.5 \text{ TFLOP/s}$

$$\text{GFLOP} \rightarrow \quad \times 10^9$$

- Frontier supercomputer: 1.7 EFLOP/s

$$\text{EFLOP} \rightarrow \quad \times 10^{18}$$

Measuring compute performance: FLOP/s

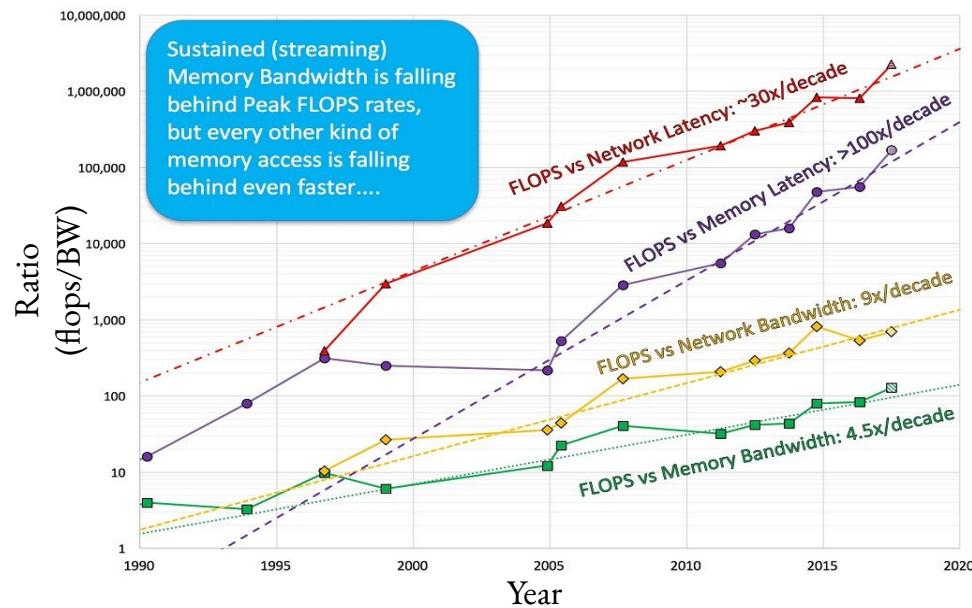


Machine balance

- Machine balance

$$(MB) = \frac{FLOP/s}{BW}$$

- Much more expensive to move data than to compute on existing data
- Ratio is worsening over the years



[McCalpin , HPC system machine balance trends, 2016]

Application FLOP/s and Bandwidth

- Application performance generally measured also with achieved FLOP/s.
- Can be computed from work complexity, $W_{app}(n)$, memory complexity $M_{app}(n)$, and application runtime, $t_{app}(n)$

$$\mathcal{P}_{app}(n) = \frac{W_{app}(n)}{t_{app}(n)}$$

Achieved performance
estimate

$$\mathcal{BW}_{app}(n) = \frac{M_{app}(n)}{t_{app}(n)}$$

Achieved bandwidth
estimate

- Can be misleading:
 - Ignores cache effects
 - High $\mathcal{P}_{app}(n)$ does not necessarily indicate optimal algorithm. Why ?

Application FLOP/s and Bandwidth

- Application performance generally measured also with achieved FLOP/s.
- Can be computed from work complexity, $W_{app}(n)$, memory complexity $M_{app}(n)$, and application runtime, $t_{app}(n)$

$$\mathcal{P}_{app}(n) = \frac{W_{app}(n)}{t_{app}(n)}$$

Achieved performance
estimate

$$\mathcal{BW}_{app}(n) = \frac{M_{app}(n)}{t_{app}(n)}$$

Achieved bandwidth
estimate

- Can be misleading:
 - Ignores cache effects
 - High $\mathcal{P}_{app}(n)$ does not necessarily indicate optimal algorithm. Why ? (No indication of useful FLOPs)
- Nevertheless, a good starting point to analyze performance of your application.

Arithmetic Intensity

- A rule of thumb for measuring ratio of compute intensity to memory intensity
- Example: (Vector addition): $u + v$

$$\left\| \begin{array}{c} | \\ = \\ | \end{array} \right\| \left\| \begin{array}{c} | \\ + \\ | \end{array} \right\|$$

$$AI = \frac{FLOP\ count}{Number\ of\ byte\ accesses}$$

$$AI = \frac{n}{3n} = \frac{1}{3} = \mathcal{O}(1)$$

Arithmetic Intensity

- A rule of thumb for measuring ratio of compute intensity to memory intensity
- Example: (Vector addition): $u + v$
- Example: (Matrix-vector product): $x = Ab$

$$AI = \frac{FLOP\ count}{Number\ of\ byte\ accesses}$$

$$AI = \frac{n}{3n} = \frac{1}{3} = \mathcal{O}(1)$$

$$AI = ? \quad \begin{array}{c} | \\ \equiv \\ | \end{array} \quad \begin{array}{c} | \\ \times \\ | \end{array}$$

Arithmetic Intensity

- A rule of thumb for measuring ratio of compute intensity to memory intensity
- Example: (Vector addition): $u + v$
- Example: (Matrix-vector product): $x = Ab$

$$AI = \frac{FLOP\ count}{Number\ of\ byte\ accesses}$$

$$AI = \frac{n}{3n} = \frac{1}{3} = \mathcal{O}(1)$$

$$AI = \frac{n^2 + n}{2n^2 + n} \cong \frac{1}{2} = \mathcal{O}(1)$$

Arithmetic Intensity

- A rule of thumb for measuring ratio of compute intensity to memory intensity
- Example: (Vector addition): $u + v$
- Example: (Matrix-vector product): $x = Ab$
- Example: (Matrix-matrix product): $C = AB$

$$AI = \frac{FLOP\ count}{Number\ of\ byte\ accesses}$$

$$AI = \frac{n}{3n} = \frac{1}{3} = \mathcal{O}(1)$$

$$AI = \frac{n^2 + n}{2n^2 + n} \cong \frac{1}{2} = \mathcal{O}(1)$$

$$AI = ? \quad \begin{array}{|c|} \hline \text{ } \\ \hline \end{array} = \begin{array}{|c|} \hline \text{ } \\ \hline \end{array} \otimes \begin{array}{|c|} \hline \text{ } \\ \hline \end{array}$$

Arithmetic Intensity

- A rule of thumb for measuring ratio of compute intensity to memory intensity
- Example: (Vector addition): $u + v$
- Example: (Matrix-vector product): $x = Ab$
- Example: (Matrix-matrix product): $C = AB$

$$AI = \frac{FLOP\ count}{Number\ of\ byte\ accesses}$$

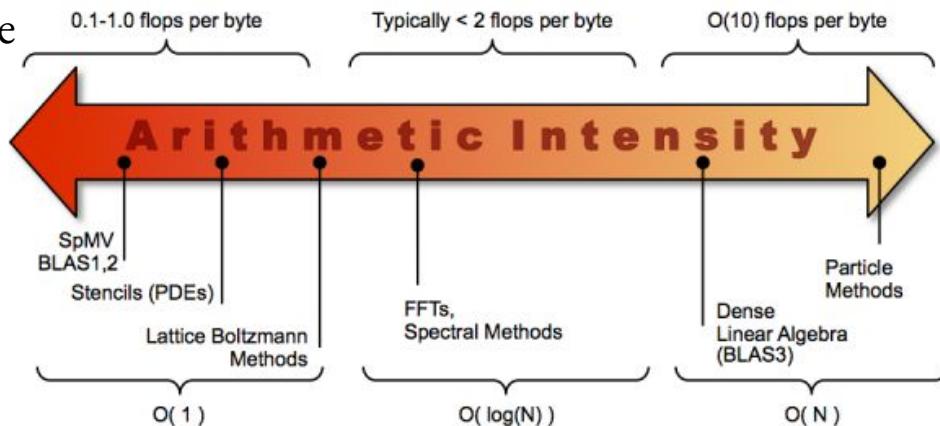
$$AI = \frac{n}{3n} = \frac{1}{3} = \mathcal{O}(1)$$

$$AI = \frac{n^2 + n}{2n^2 + n} \cong \frac{1}{2} = \mathcal{O}(1)$$

$$AI = \frac{2n^3}{3n^2} = \frac{2n}{3} = \mathcal{O}(n)$$

Arithmetic Intensity

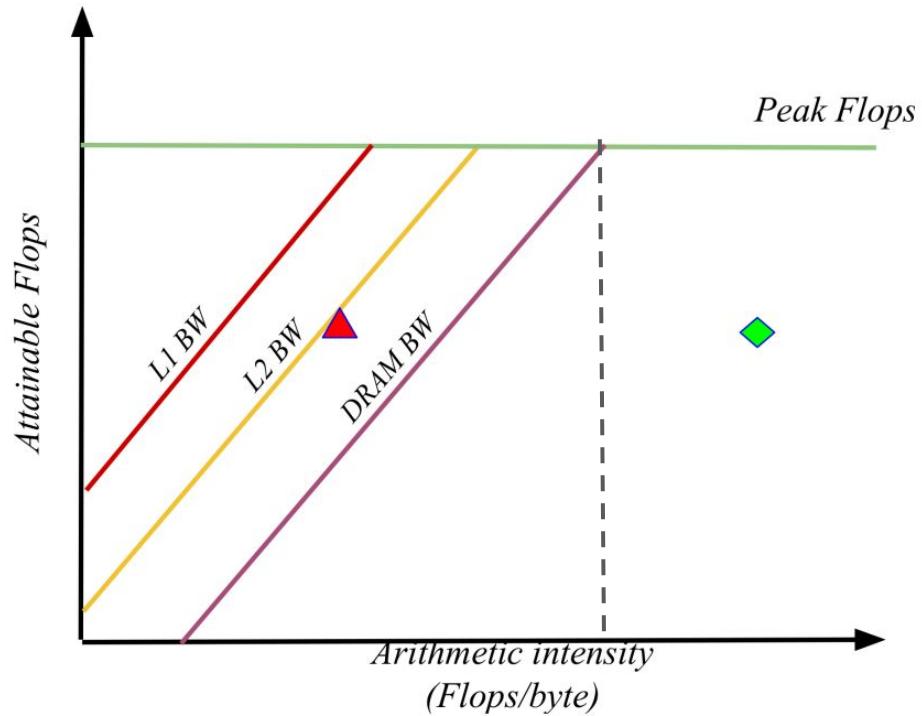
- Due to the imbalance in the machine, higher the arithmetic intensity of an operation, the better.
- Reframe algorithms into building blocks consisting of higher arithmetic intensity operations
- Particle methods, BLAS 3 operations can be very performant.



[Roofline model, crd.lbl.gov]

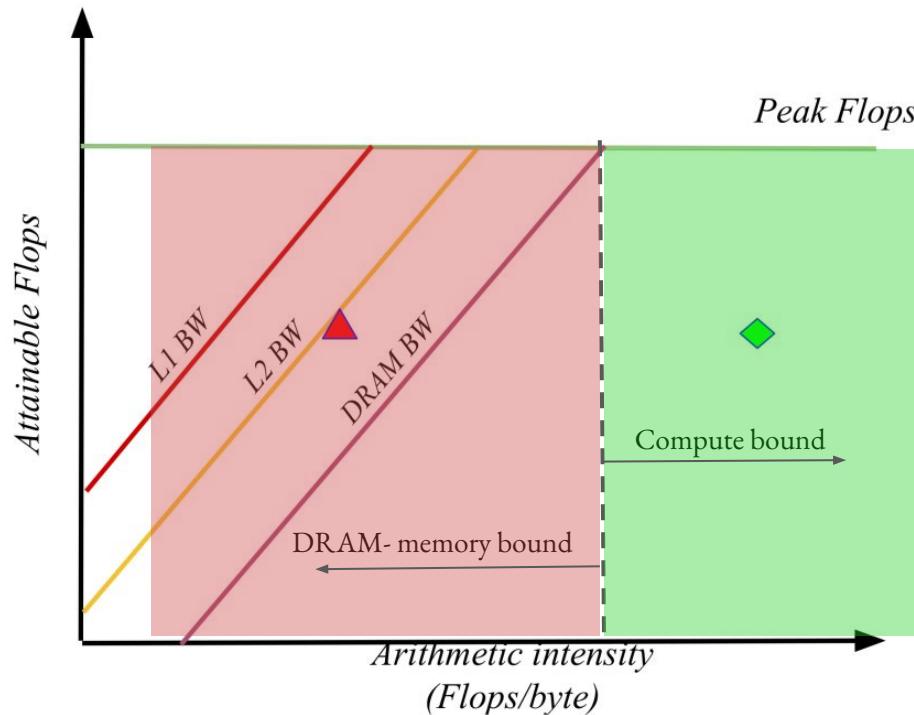
Measuring algorithm performance: Roofline

- A throughput oriented performance model
- Independent of ISA and architecture
- Components:
 - Machine parameters
 - Theoretical application bounds
 - Actual application execution
- Estimate whether an application is memory-bound or compute-bound



Measuring algorithm performance: Roofline

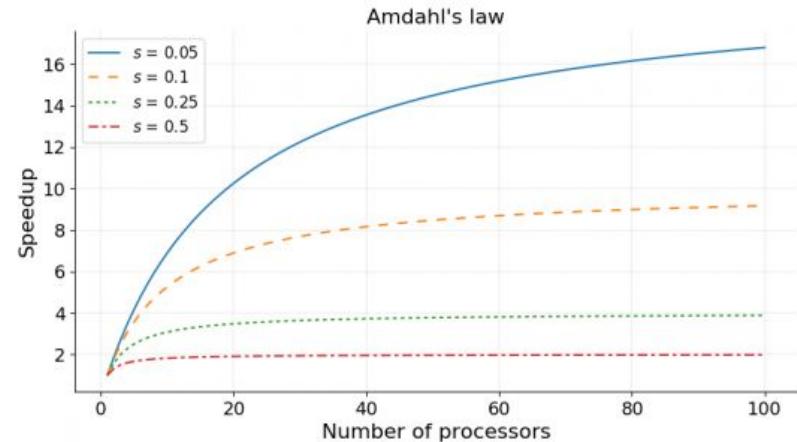
- A throughput oriented performance model
- Independent of ISA and architecture
- Components:
 - Machine parameters
 - Theoretical application bounds
 - Actual application execution
- Estimate whether an application is memory-bound or compute-bound



Measuring performance scaling: Strong scaling

- Speedup of some application: $S = \frac{t_1}{t_N}$
- Consider an application with p parallelizable component and s serial component, $s = 1 - p$
- Effective speedup:

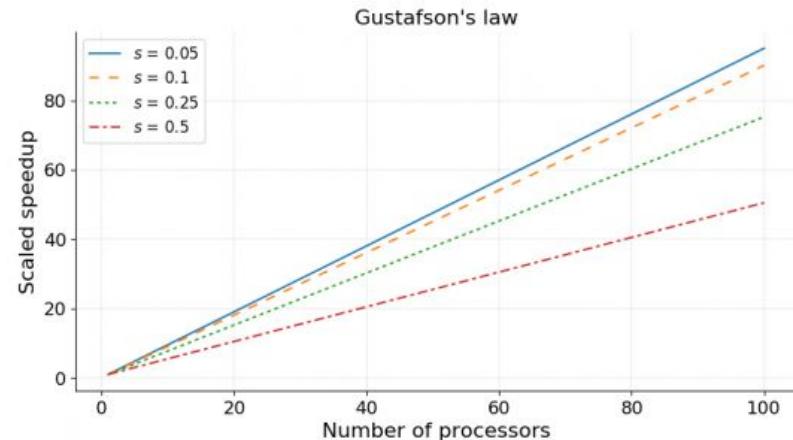
$$S_{strong} = \frac{1}{s + \frac{p}{N}}$$



- Serial components kill performance

Measuring performance scaling: Weak scaling

- Strong scaling keeps problem size same, increase number of available resources
- More realistic:
 - Increase problem size with available resources.
 - Aim to saturate each available resource.



$$S_{weak} = s + p \times N$$

Summary and takeaways

- To parallelize your code, look at:
 - Dependencies
 - Available resources
- Decompose your problem to maximize parallelism
- Reframe your algorithm to have building blocks of high arithmetic intensity
- Be aware of cache effects and NUMA effects
- Code → measure performance → re-code → measure

Next week

- Parallel programming models
 - Shared and distributed parallel programming models
 - CUDA (NVIDIA), HIP (AMD), SYCL (Intel), Metal (Apple)
- GPU hardware architecture
 - Important differences to CPUs
 - Compute and Memory hierarchies
 - Brief look at ISA